# FORGE

## Framework for Orchestrated Requirements, Governance & Engineering

A structured software development system for OpenCode

# What is FORGE?

**FORGE** is a comprehensive methodology for AI-assisted software development that combines:

- **Structured progressive context** that prevents inconsistent decisions
- **Adaptive process** that calibrates ceremony to complexity
- **Persistent knowledge** that survives session boundaries

> Built natively for OpenCode, synthesising the best of BMAD Method and Speckit

# The Problem

## 5 Critical Problems in AI-Assisted Development

1. **Context Gap** - Agents in different sessions make conflicting decisions
2. **Ceremony Trap** - Too much process kills speed, too little kills quality
3. **Knowledge Evaporation** - Decisions disappear at the end of the session
4. **Consistency Entropy** - Teams of 15+ developers produce inconsistent code
5. **Illusion of Productivity** - Speed without direction is waste

# The FORGE Solution

## 6 Core Principles

| Principle | Description |
|---|---|
| **Progressive Context Engineering** | Each phase produces a document that becomes context for the next phase |
| **Constitutional Governance** | Immutable principles govern every decision |
| **Adaptive Ceremony** | Process depth scales with complexity |
| **Adversarial Quality** | Reviews MUST find problems |
| **Persistent Knowledge** | Decisions and lessons survive across sessions |
| **Bidirectional Traceability** | Requirements → code → tests (and back) |

# System Architecture

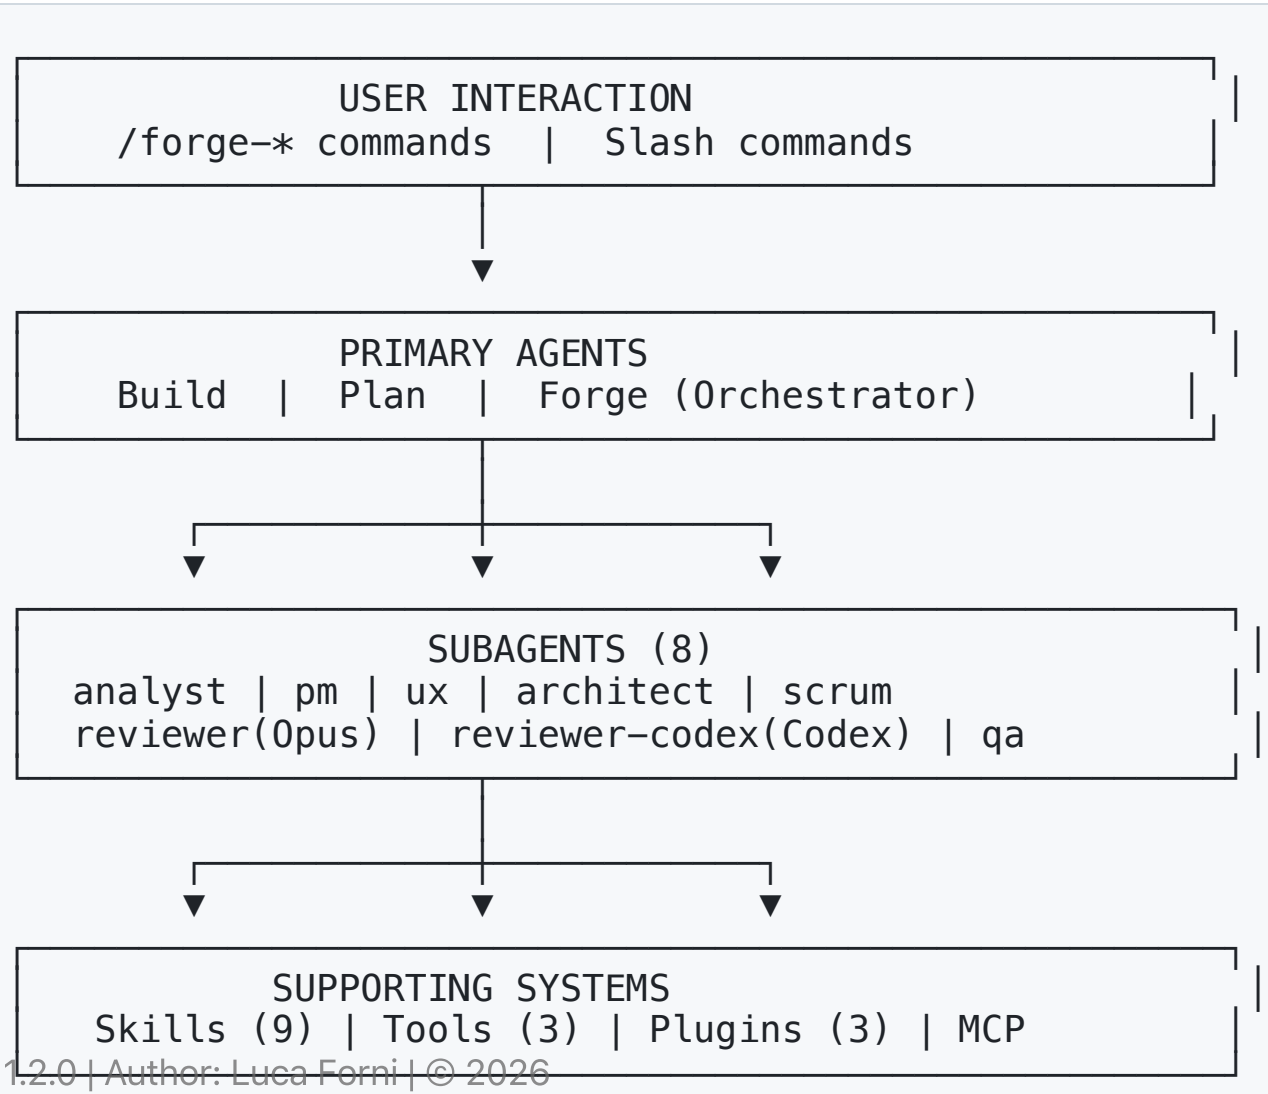

# 5 Workflow Tracks

```
Complexity ──────────────────────────────────────────────────────> High

Hotfix      Quick       Feature      Epic         Product
──────      ──────      ──────────   ──────       ──────────
1 file      1–5 tasks   5–20 tasks   20–50+       New product
< 30 min    < 1 day     1–5 days     1–4 weeks    4+ weeks
No docs     Tech spec   Spec+Plan    Full chain   Full chain
                                     + Sprint     + Constitution
```

**The process automatically adapts to task complexity**

# Track: Hotfix

**When:** Critical bug, 1-2 files, < 30 minutes

```
/forge-hotfix "Login endpoint returns 500 when user has no profile picture"
```

**Workflow:**

1. ✓ Diagnose - Identify root cause
2. ✓ Fix - Apply minimal, targeted fix
3. ✓ Verify - Run existing tests
4. ✓ Review - Quick self-review vs constitution

**Output:** Structured commit message (no additional documents)

# Track: Quick

**When:** Small feature, 1-5 tasks, < 1 day

```
/forge-quick "Add forgot password feature with 1-hour expiry token"
```

**Workflow:**

1. ✓ Quick Spec - Conversation → `tech-spec.md`
2. ✓ Implement - Implement task by task
3. ✓ Test - Generate unit tests
4. ✓ Review - Adversarial self-review

**Output:** `.forge/specs/NNN-name/tech-spec.md` + code + tests

# Track: Feature

**When:** Medium feature, 5-20 tasks, 1-5 days

```
/forge-specify "Add OAuth2 authentication with Google and GitHub"
/forge-clarify    # Resolve ambiguities
/forge-ux         # UX design: personas, wireframes, accessibility
/forge-plan       # Technical plan
/forge-analyze    # Cross-validate spec vs plan
/forge-tasks      # Task breakdown with dependencies
/forge-implement  # Implement
/forge-review     # Adversarial review (7 dimensions, dual-model: Claude Opus + GPT-Codex)
```

**Output:** `spec.md`, `design-spec.md`, `user-journey.md`, `plan.md`, `tasks.md`, optional ADRs

# UX Phase: `/forge-ux`

**When:** Feature with user interface (web, mobile, design system)

```
/forge-ux "Login page with OAuth"
```

| Artifact | Content | Focus |
|---|---|---|
| `design-spec.md` | ASCII wireframes, UI components, interaction specs | 🖼️ wireframes · 🎨 design system · 📱 responsive |
| `user-journey.md` | Personas, scenarios, user flows | 🗺️ journeys · ♿ WCAG 2.1 AA |

# Review: 7 Quality Dimensions

## The review covers all 7 dimensions:

```
Correctness           → logic, edge cases
Security              → vulnerabilities, injection
Performance           → slow queries, memory leaks
Maintainability       → readability, coupling
Constitution          → adherence to principles
Test-Spec Coherence   → tests cover spec requirements
UX Quality            → accessibility, usability, design consistency
```

## UX issue examples:

- `[HIGH]` Button without `aria-label` — not accessible with screen reader
- `[MEDIUM]` Form without focus management — degraded UX after submit
- `[LOW]` Colour contrast 2.8:1 — below WCAG AA threshold (4.5:1)

# Track: Epic

**When:** Complex feature set, 20-50+ tasks, 1-4 weeks

```
/forge-brief          # Product brief
/forge-prd            # Full PRD
/forge-architecture   # Architecture + ADR
/forge-sprint         # Sprint planning
/forge-story          # Prepare story
/forge-implement      # Implement story
/forge-review         # Dual review (AI + Human)
/forge-retro          # Retrospective
```

**Output:** Brief, PRD, Architecture, Epic, Stories, Sprint Status, ADR

# Track: Product

**When:** New product, greenfield, 4+ weeks

```
/forge-init          # Setup + Constitution
```

Then follows the Epic workflow with the addition of:

- **Constitution** - Immutable governance document
- **UX Phase** - `/forge-ux` for products with user interface

**Full output:** All Epic track documents + Constitution

# The 8 Specialized Subagents

| Agent | Model | Role |
|---|---|---|
| **forge-analyst** | Sonnet 4.5 | Exploration, research, scope detection |
| **forge-pm** | Opus 4.6 | Requirements, spec, PRD, user stories |
| **forge-ux** | Opus 4.6 | User journeys, wireframes, accessibility, design spec |
| **forge-architect** | Opus 4.6 | Architecture, ADR, technical planning |
| **forge-scrum** | Sonnet 4.5 | Sprint planning, story management |
| **forge-reviewer** | Opus 4.6 | Adversarial review — Task A (7 dimensions) |
| **forge-reviewer-codex** | GPT-Codex | Adversarial review — Task B, independent |
| **forge-qa** | Sonnet 4.5 | Test strategy, test generation |

> `/forge-review` launches **forge-reviewer + forge-reviewer-codex in parallel** and synthesises findings — *consensus findings* (both models) carry the highest priority

# Progressive Context Chain

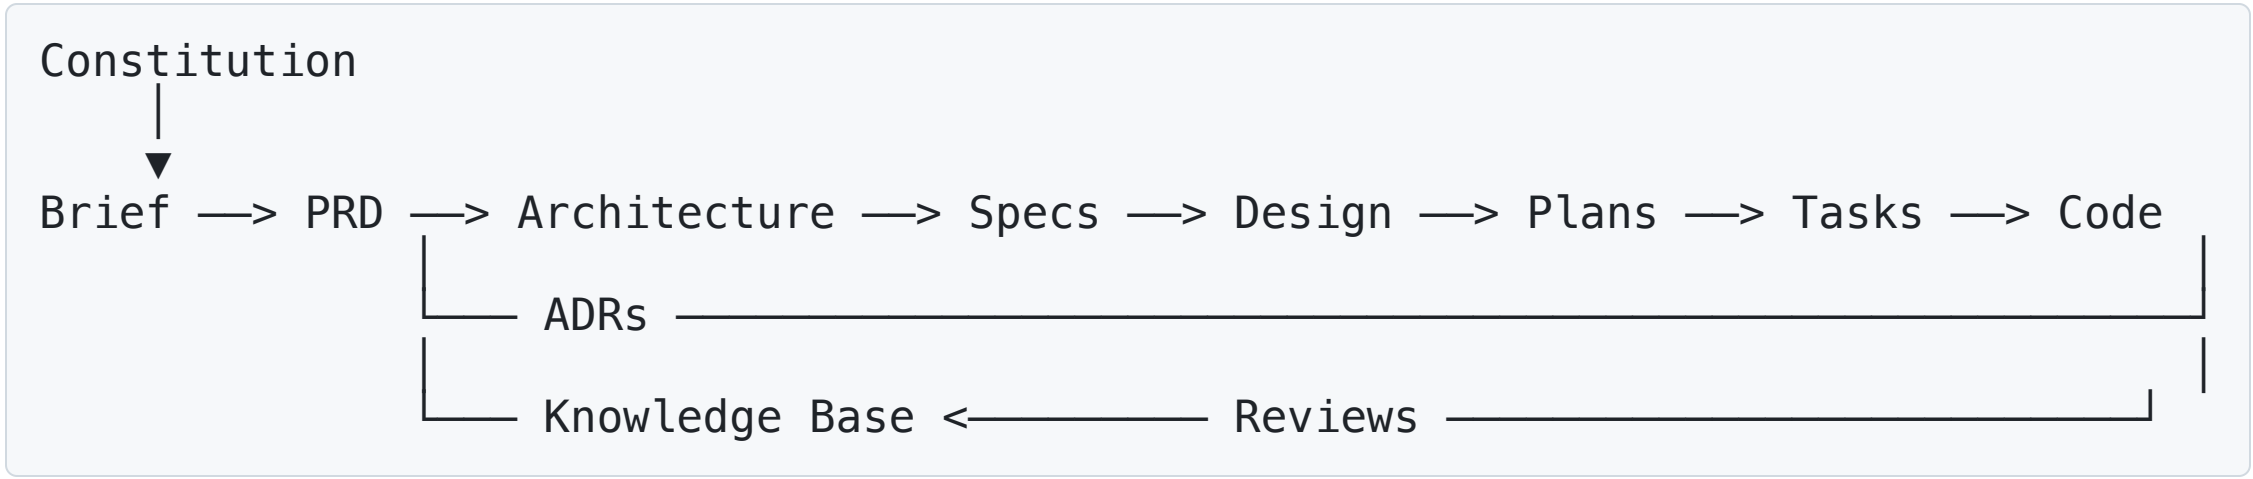


**Each phase receives structured context from the previous phase**

- Prevents inconsistent decisions
- Guarantees architectural alignment
- Accumulates knowledge over time

# Constitutional Governance

The **Constitution** is the highest-authority governance document

## 9 Standard Articles

1. **Core Principles** - Non-negotiable principles
2. **Technology Stack** - Approved tech stack
3. **Architecture Patterns** - Mandatory patterns
4. **Quality Standards** - Quality thresholds
5. **Security** - Security requirements
6. **Error Handling** - Error handling standards
7. **Naming & Conventions** - Naming conventions
8. **Testing Standards** - Required tests
9. **Operational Requirements** - Operational standards

# Persistent Knowledge Base

## 3 Types of Artefacts

```
.forge/knowledge/
├── adr/                        # Architecture Decision Records
│   ├── 001-database-choice.md
│   └── 002-auth-strategy.md
├── decision-log.md             # Session-extracted decisions
└── lessons-learned.md          # Post-mortem insights
```

**Features:**

- ✓ Created automatically by the `session-knowledge` plugin
- ✓ Persist across sessions
- ✓ Loaded as context every session
- ✓ Prevent repeated mistakes

# Adversarial Review

## Two Models in Parallel — Reviews that MUST Find Problems

`forge-reviewer` (Claude Opus) + `forge-reviewer-codex` (GPT-Codex) examine across **7 dimensions:**

1. **Correctness** - Logic, edge cases, errors
2. **Security** - Vulnerabilities, input validation
3. **Performance** - Slow queries, memory leaks
4. **Maintainability** - Readability, coupling
5. **Constitution Compliance** - Adherence to principles
6. **Test-Spec Coherence** - Tests cover spec requirements
7. **UX Quality** - Accessibility (WCAG 2.1 AA), usability, design consistency

**Minimum 5 combined issues** across both models — *consensus findings* (flagged by both)

# Dual Review Process

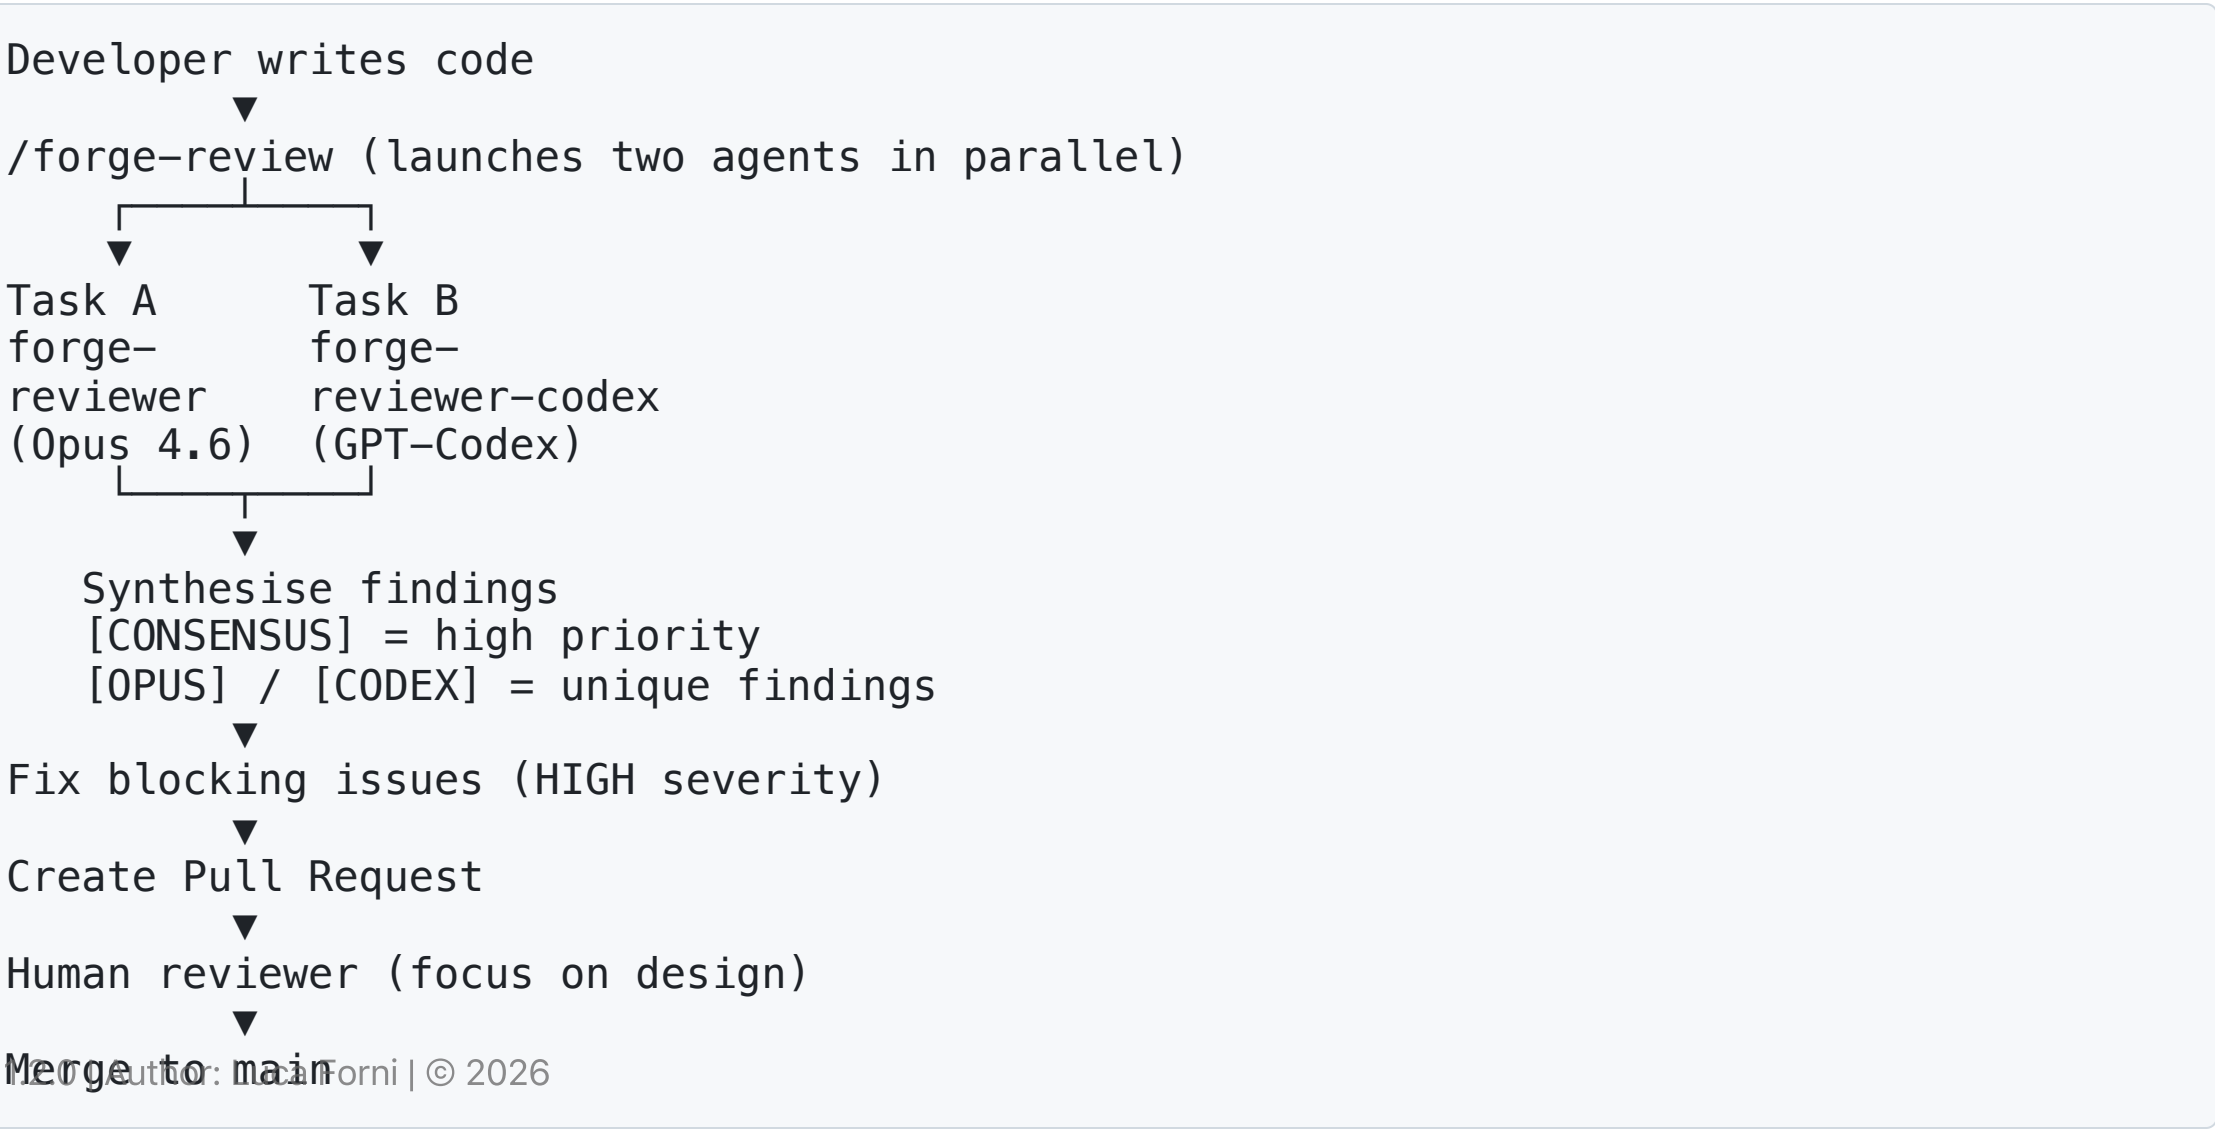

# Bidirectional Traceability

```
Requirement FR-001 (spec.md)
  ↓
Technical approach (plan.md, section 3.2)
  ↓
Task 2.1 (tasks.md)
  ↓
Source file (src/auth/login.ts)
  ↓
Test file (src/auth/__tests__/login.test.ts)
```

**Custom tool:** `/forge-analyze` generates traceability matrix

- Identifies unimplemented requirements
- Identifies code without requirements (orphan code)
- Verifies test coverage per requirement

# Brownfield Support

## Onboarding Existing Codebases

```
/forge-init  # In existing project
```

### FORGE analyses:

- Project structure and languages
- Detected frameworks and architectural patterns
- Naming and organisation conventions
- Potential tech debt

### Output:

- Constitution auto-generated based on existing code
- AGENTS.md derived from `.eslintrc`, `tsconfig.json`, detected patterns

# 9 Dynamic Skills

| Skill | Used by | Purpose |
|---|---|---|
| `adversarial-review` | forge-reviewer | Mandatory review protocol |
| `advanced-elicitation` | pm, architect | Deep analysis techniques |
| `scope-detection` | Forge orchestrator | Assess complexity, recommend track |
| `test-strategy` | forge-qa, Build | Adaptive test strategy per track |
| `brownfield-analysis` | forge-analyst | Existing codebase analysis |
| `constitution-compliance` | architect, reviewer | Verify constitution compliance |
| `context-chain` | All agents | Load correct upstream documents |
| `ux-design` | forge-ux | Generate user journeys, wireframes, a11y specs |
| `ux-review` | forge-reviewer | 7th review dimension: UX quality & accessibility |

**Loaded on-demand to save context window**

# Automatic Plugins (1/2)

## session-knowledge

- Automatically extracts decisions and lessons
- Appends to `decision-log.md` and `lessons-learned.md`
- Injects knowledge during context compaction

## pre-commit-gate

- Validates spec-code consistency before commit
- Verifies completed tasks and existing tests
- Advisory (does not block, but warns)

# Automatic Plugins (2/2)

## spec-watcher

- Monitors changes to `.forge/specs/`
- Detects inconsistencies with plan/tasks
- Suggests `/forge-analyze`

**Plugins automate governance and knowledge management without manual intervention**

# Benefits for Developers

| Benefit | How FORGE Ensures It |
|---|---|
| **Fast onboarding** | Spec, plan, ADR provide full context |
| **Quality on the first try** | Structured context → better AI decisions |
| **Less rework after review** | AI review catches issues before human review |
| **Clear requirements** | `/forge-clarify` surfaces ambiguities before implementation |
| **Confidence in decisions** | Constitution and ADR validate choices |

# Benefits for Teams

| Benefit | How FORGE Ensures It |
|---|---|
| **Consistent code** | Constitution enforces patterns; all agents follow the same rules |
| **Reduces architecture drift** | ADRs prevent contradictory decisions |
| **Efficient sprint management** | `/forge-sprint` and `/forge-status` automate overhead |
| **Effective retrospectives** | `/forge-retro` produces actionable insights |
| **Faster reviews** | AI handles mechanical checks, humans focus on design |
| **Smooth onboarding** | New members read `.forge/` to understand the whole project |
| **Safe parallel development** | Architecture defines boundaries; ADRs prevent conflicts |

# Benefits for Enterprise

| Benefit | How FORGE Ensures It |
|---|---|
| **Audit trail compliance** | Constitution + ADR + decision log = complete rationale |
| **Risk management** | Spec includes risk sections; adversarial review finds issues early |
| **Knowledge retention** | When developers leave, decisions remain |
| **Process standardisation** | All teams use the same workflows, templates, protocols |
| **Tech debt visibility** | Brownfield analysis + traceability expose gaps |
| **Quality metrics** | Sprint velocity, issue count, spec completeness |
| **Governance without bottleneck** | Constitution + skills enforce standards automatically |

# FORGE vs Alternatives

## vs No Methodology ("Vibe Coding")

| | No Methodology | FORGE |
|---|---|---|
| Initial speed | Very fast | Moderate (planning overhead) |
| Speed over time | Slows down (debt accumulates) | Sustained (compounding knowledge) |
| Consistency | Random | Enforced |
| Rework | High (30-50%) | Low (< 15%) |
| Onboarding | Weeks | Days |

# FORGE vs BMAD vs Speckit

| | BMAD | Speckit | FORGE |
|---|---|---|---|
| Agent system | Personas (1 LLM) | None | Real subagents |
| Tracks | 3 | 1 | 5 |
| Governance | Weak | Constitution | Constitution + ADR + KB |
| Knowledge persistence | None | Per-branch | Cross-session KB |
| Review | Adversarial | None | Dual Adversarial |
| Brownfield | Limited | Limited | Structured |
| Platform | IDE-agnostic | Agent-agnostic | OpenCode-native |

# When to Use FORGE

## ✅ Ideal Use Cases

| Scenario | Track | Why |
|---|---|---|
| New SaaS from scratch | Product | Prevents tech debt |
| Enterprise feature | Epic | Prevents scope creep |
| Large codebase | Feature/Quick | Maintains consistency |
| Compliance-heavy | Epic/Product | Complete audit trail |
| Team with turnover | Any | Retains knowledge |

# When NOT to Use FORGE

## ❌ Overkill Scenarios

| Scenario | Better Alternative |
|---|---|
| One-off scripts or utilities | Write code directly |
| Tutorial / learning projects | Focus on learning, not process |
| Hackathon prototype (discard after) | Speed > structure |
| Projects with < 1 week total lifespan | The docs outlive the code |
| Safety-critical systems | FORGE + formal verification |
| Multi-team platform (50+ devs) | FORGE + enterprise program management |

# The Cost of Structure

## Realistic Overhead

| Activity | Overhead |
|---|---|
| Constitution | 1-2 hours (once) |
| Spec (Feature) | 20-30 min |
| Plan | 15-20 min |
| Analyze | 5 min |
| Review | 10-15 min |
| Retro | 15-20 min/sprint |

**Total per feature: ~1-1.5 hours**

# Break-Even Analysis

```
Without FORGE:
  Implementation: 3 days
  Rework after review: 0.5 days (frequent)
  Rework after production: 1 day (occasional)
  Total: 3.8 days

With FORGE:
  Planning + review: 0.2 days
  Implementation: 2.5 days (better context = faster)
  Rework after review: 0.1 days (AI review catches issues)
  Rework after production: 0.1 days (rare with dual review)
  Total: 2.9 days
```

**Break-even typically at the 2nd–3rd feature of the project**

# Compound Effect

```
Session 1:    You pay the constitution cost → high overhead, zero benefit
Session 10:   Constitution prevents a bad decision → benefit > costs
Session 50:   New dev productive in hours thanks to .forge/
Session 200:  Auditor asks encryption rationale → ADR-012 → 1 day vs 1 week
```

**Structure has decreasing costs and compounding returns**

The question is not "can we afford the overhead?" but

**"can we afford NOT to have it?"**

# Technical Components

## 10 Agents | 21 Commands | 9 Skills

```
.opencode/
├── agents/          # 8 specialized subagents + forge orchestrator + Build + Plan
├── commands/        # 21 slash commands
├── skills/          # 9 dynamic skills
├── tools/           # 3 custom tools
├── plugins/         # 3 automation plugins
└── templates/       # 11 document templates
```

**All configured in `opencode.json`**

# Quick Start

## 1. Installation

```
# Copy .opencode/ into your project
cp -r path/to/forge/.opencode/ your-project/.opencode/
cp path/to/forge/opencode.json your-project/opencode.json
```

## 2. Verify

```
cd your-project
opencode
> /forge-help
```

## 3. First Workflow (Quick)

```
> /forge-quick "Add health check endpoint returning app version"
```

**Output:** Spec + code + tests in < 5 minutes

# Directory Structure

```
.forge/
├── constitution.md                # Governance
├── product/
│   ├── brief.md                   # Strategic vision
│   └── prd.md                     # Full requirements
├── ux/
│   └── design-system.md           # Shared tokens & components
├── architecture/
│   └── architecture.md            # Technical design
├── specs/
│   └── NNN-name/                  # Per-feature specs
│       ├── spec.md
│       ├── design-spec.md         # UX/UI design (wireframes, a11y)
│       ├── user-journey.md        # Personas & user journeys
│       ├── plan.md
│       └── tasks.md
├── epics/                         # Epic breakdown
├── sprints/                       # Sprint tracking
└── knowledge/                     # Persistent memory
    ├── adr/                       # Decision records
    ├── decision-log.md
    └── lessons-learned.md
```

# Model Strategy

## Differentiation by Cognitive Demand

| Model | When | Agents |
|---|---|---|
| Claude Opus 4.6 | Deep reasoning, architectural decisions, adversarial review | forge-pm, forge-architect, forge-reviewer |
| Claude Sonnet 4.5 | Speed, good-enough reasoning, analysis, sprint mgmt | Forge, forge-analyst, forge-scrum, forge-qa, Build, Plan |
| GPT-5.2-Codex | Independent adversarial review (second model) | forge-reviewer-codex |

## Models provided via GitHub Copilot subscription

# Incremental Adoption Path

## No Need to Adopt Everything at Once

**Week 1-2:** Hotfix + Quick only

- Get familiar with the workflow without changing existing process

**Week 3-4:** Feature track for a medium feature

- Experience the full spec-plan-implement-review cycle

**Month 2:** Epic track for a larger initiative

- Add sprint management

**Month 3+:** Evaluate Product track and constitutional governance

# Essential Commands

## Quick Reference Card

```
Tracks:           /forge-hotfix   Bug fix, 1 file, < 30 min
                  /forge-quick    Small feature, 1-5 tasks
                  /forge-specify  Medium feature (start here)
                  /forge-brief    Large epic
                  /forge-init     New product

Feature Flow:     specify → clarify → ux → plan → analyze
                  → tasks → implement → review

UX Commands:      /forge-ux           User journeys, wireframes, a11y
                  /forge-wireframe    ASCII wireframe generation

Status:           /forge-status   Sprint dashboard
                  /forge-help     Context-aware help

Knowledge:        /forge-adr      Create ADR
                  /forge-retro    Sprint retrospective
```

# Use Case: OAuth Feature

```
# 1. Specify
/forge-specify "Add OAuth2 auth with Google and GitHub"

# 2. Clarify ambiguities
/forge-clarify

# 3. UX Design
/forge-ux  # login wireframe, user journey, a11y

# 4. Technical plan
/forge-plan

# 5. Validate consistency
/forge-analyze

# 6. Task breakdown
/forge-tasks

# 7. Implement
/forge-implement

# 8. Adversarial review (7 dimensions, dual-model: Claude Opus + GPT-Codex)
/forge-review
```

# Team Workflow

## Multi-Developer Feature Development

```
Developer A: Epic 1, Stories S001–S004 (Core Payments)
Developer B: Epic 2, Stories S001–S003 (Subscriptions)
Developer C: Epic 3, Stories S001–S003 (Webhooks)
```

**Shared artefacts** (via git):

- `.forge/constitution.md` - Everyone follows the same principles
- `.forge/architecture/` - Consistent technical decisions
- `.forge/knowledge/adr/` - Everyone sees all decisions
- `.forge/sprints/` - Scrum master updates centrally

**Conflict prevention:**

- Architecture defines module boundaries

# Sprint Ceremonies with FORGE (1/2)

## Sprint Planning

```
/forge-sprint
# → Review velocity
# → Select stories from backlog
# → Assign to developer
# → Update sprint-status.yaml
```

## Daily Standup

```
/forge-status
# → See assigned stories and status
# → Flag blockers
```

# Sprint Ceremonies with FORGE (2/2)

## Sprint Retrospective

```
/forge-retro
# → Automatic velocity calculation
# → Lessons learned → knowledge base
```

# Knowledge Base Maintenance

| Task | Frequency | Owner |
|---|---|---|
| Review decision-log, promote to ADR | Weekly | Tech Lead |
| Archive stale lessons-learned | Per sprint | Scrum Master |
| Review ADR staleness | Monthly | Architect |
| Verify constitution accuracy | Quarterly | Tech Lead + PM |

**15 minutes/week for KB maintenance**

# Customisation

## FORGE is Fully Customisable

- **Constitution Template** - Adapt articles to your domain
- **Agent System Prompts** - Tune for team style
- **Skills** - Add domain-specific reasoning
- **Templates** - Modify document structure
- **Commands** - Create custom workflows
- **Plugins** - Project-specific automation

**See:** `.opencode/docs/FORGE-CUSTOMIZATION.md`

# CI/CD Integration

## Pre-commit Gate Plugin

```
// Runs before commit
- Check tasks completed in tasks.md
- Verify tests exist for modified files
- Validate [NEEDS CLARIFICATION] resolved
- Check constitution compliance verified
```

**Non-blocking** (advisory), but provides a clear signal

# Advanced Elicitation Techniques

## 6 Deep Analysis Techniques

1. **Pre-mortem Analysis** - Imagine the feature failed — what went wrong?
2. **First Principles Thinking** - Decompose into fundamental components
3. **Red Team / Blue Team** - Attacker vs defender
4. **Socratic Questioning** - Deep questions about assumptions
5. **Constraint Removal** - What if we had no constraints?
6. **Inversion Analysis** - How would we guarantee it FAILS?

**Used by `forge-pm` and `forge-architect` for spec and architecture analysis**

# Scope Detection

## Automatic Complexity Assessment

**7 Factors:**

- Files affected (1-2 → 50+)
- Estimated tasks (1 → 50+)
- New dependencies (0 → Stack decision)
- Schema changes (None → Full design)
- API surface changes
- Cross-module impact
- Need for new patterns

**Output:** Structured JSON with recommended track + reasoning

**User always has the final say**

# Adaptive Test Strategy

## Coverage Requirements per Track

| Track | Required Tests |
|---|---|
| **Hotfix** | Regression test for the bug only |
| **Quick** | Unit tests for new/modified code |
| **Feature** | Unit + Integration tests |
| **Epic** | Unit + Integration + E2E |
| **Product** | Unit + Integration + E2E + Performance benchmarks |

**Defined by the `test-strategy` skill, executed by `forge-qa`**

# ADR (Architecture Decision Records)

## When to Create an ADR

- Database, framework, major library choice
- API style decision (REST vs GraphQL vs gRPC)
- Defining pattern (event-driven vs synchronous)
- Trade-off (consistency vs availability)
- Any decision someone will question later

## Format

```
# Context: Why is this decision needed?
# Options: What alternatives were considered?
# Decision: What was chosen and why?
# Consequences: Positive, negative, neutral effects
# Constitution Alignment: Which articles does it support?
```

# Retrospectives

## `/forge-retro` Output

```
Sprint 2 Retrospective
======================
Velocity: 32 pts (planned: 34, previous: 28)

What went well:
  - Stripe integration straightforward thanks to clear ADRs
  - Task parallelism markers saved time

What could improve:
  - Webhook testing required manual Stripe CLI setup (not in spec)
  - OAuth PR took 2 days human review (bottleneck)

Action items:
  - Add webhook test setup instructions to spec template
  - Rotate PR reviewer to avoid single-person bottleneck

Lessons → .forge/knowledge/lessons-learned.md
```

# Traceability Matrix Example

```
Requirement FR-001 (Login with OAuth)
  → Plan Section 3.2 (OAuth flow implementation)
    → Task 2.1 (Implement OAuth strategy)
      → src/auth/oauth-strategy.ts ✓
        → src/auth/__tests__/oauth-strategy.test.ts ✓

Requirement FR-002 (Account linking)
  → Plan Section 3.3 (Link accounts)
    → Task 3.1 (Link endpoint) → [NOT IMPLEMENTED] ⚠️

Requirement NFR-001 (P95 < 200ms)
  → [NO TASK] → [NOT IMPLEMENTED] ⚠️
```

**Generated by the `trace-requirements` tool**

# Success Metrics

## How to Measure FORGE Success

| Metric | Target | How Measured |
|---|---|---|
| Rework rate | < 15% | Issues after PR merge / Total PRs |
| Onboarding time | < 3 days | New dev to first meaningful commit |
| Architecture drift incidents | 0 | ADR violations detected |
| Knowledge retention | 100% | Documented decisions / Total decisions |
| Sprint velocity | +20% after 3 sprints | Story points delivered |
| Code consistency | > 90% | Constitution compliance score |

# Resources & Documentation

## Full Documentation in `.opencode/docs/`

- **FORGE-GUIDE.md** - Complete usage guide
- **FORGE-PHILOSOPHY.md** - Principles and benefits
- **FORGE-PROJECT-PLAN.md** - System architecture
- **FORGE-CUSTOMIZATION.md** - How to customise
- **FORGE-DECISIONS.md** - Methodology decision records

## Community & Support

- GitHub: `lucaforni/forge`
- OpenCode Docs: `https://opencode.ai/docs`

# Future Roadmap

## Possible Evolutions

- **AI Pair Programming Mode** - Forge assists in real-time during coding
- **Multi-repo Support** - FORGE coordination across microservices
- **Custom Domain Templates** - Pre-built constitutions for fintech, healthcare, gaming
- **Integration with Project Management Tools** - Jira, Linear, Azure DevOps
- **Visual Architecture Diagrams** - Auto-generate from architecture.md
- **Metrics Dashboard** - Real-time visibility on velocity, quality, debt

# Key Takeaways

## Why FORGE Changes the Game

1. **Structured context** prevents inconsistencies across AI sessions
2. **5 workflow tracks** adapt the process to complexity
3. **UX-first design** includes user journeys, wireframes and accessibility in the workflow
4. **Constitutional governance** ensures a uniform quality bar
5. **Persistent knowledge base** eliminates knowledge loss
6. **Adversarial review** catches issues before production (7 dimensions, dual-model)
7. **Native OpenCode integration** leverages the full platform
8. **Brownfield support** onboards existing codebases
9. **Team-ready** supports 15+ developers with parallel development

# Get Started Today (1/2)

## 3 Steps to Try FORGE

### 1. Setup (5 minutes)

```
cp -r .opencode/ your-project/
cd your-project && opencode
```

### 2. First Task (5 minutes)

```
> /forge-quick "Your first small feature"
```

# Get Started Today (2/2)

## 3 Steps to Try FORGE

### 3. Evaluate Results

- Is the produced spec clear?
- Is the implementation clean?
- Were tests generated correctly?

**If yes → Continue with Feature track**

**If no → Contact support for tuning**

# Questions?

## Contacts:

- Documentation: `.opencode/docs/`
- GitHub Issues: `lucaforni/forge`
- OpenCode Docs: `https://opencode.ai/docs`

# Thank You!

**FORGE**

Framework for Orchestrated Requirements, Governance & Engineering

> Enterprise software development with AI — structured and sustainable

**Author:** Luca Forni

🔗 linkedin.com/in/lucaforni

🔗 github.com/lucaforni

*Version 1.2.0 | MIT License | 2026*